

Machine Learning - Course Project

Oliver Gonzalez

November 18, 2017

Background

Using devices such as Jawbone Up, Nike FuelBand, and Fitbit it is now possible to collect a large amount of data about personal activity relatively inexpensively. These type of devices are part of the quantified self movement “ a group of enthusiasts who take measurements about themselves regularly to improve their health, to find patterns in their behavior, or because they are tech geeks. One thing that people regularly do is quantify how much of a particular activity they do, but they rarely quantify how well they do it. In this project, your goal will be to use data from accelerometers on the belt, forearm, arm, and dumbbell of 6 participants. They were asked to perform barbell lifts correctly and incorrectly in 5 different ways. More information is available from the website here: <http://groupware.les.inf.puc-rio.br/har> (see the section on the Weight Lifting Exercise Dataset).

Data Processing and Data cleaning

The data for this project come from this source: <http://groupware.les.inf.puc-rio.br/har>. If you use the document you create for this class for any purpose please cite them as they have been very generous in allowing their data to be used for this kind of assignment.

The training data for this project are available here:

<https://d396qusza40orc.cloudfront.net/predmachlearn/pml-training.csv> The test data are available here: <https://d396qusza40orc.cloudfront.net/predmachlearn/pml-testing.csv>. We must Download and clean de data from division errors and empty strings replacing them by NA values.

Cross Validation & Data Partitioning

Now lets partition de data into a training (60%) and test (40%) sets. Also lets remove any NA and DIV/0 values.

```
train_url <- "http://d396qusza40orc.cloudfront.net/predmachlearn/pml-
training.csv"
test_url <- "http://d396qusza40orc.cloudfront.net/predmachlearn/pml-
testing.csv"
training_set<- read.csv(url(train_url), na.strings=c("NA","#DIV/0!", ""))
testing_set <- read.csv(url(test_url), na.strings=c("NA","#DIV/0!", ""))
inTrain <- createDataPartition(y=training_set$classe, p=0.60, list=FALSE)
myTraining <- training_set[inTrain, ]; myTesting <- training_set[-inTrain, ]
dim(myTraining); dim(myTesting)
```

```

## [1] 11776    160
## [1] 7846    160
removeNACols  <- function(x) { x[, colSums( is.na(x) ) < nrow(x) ] }
myTraining <- removeNACols(myTraining)
myTesting  <- removeNACols(myTesting)

complete  <- function(x) {x[,sapply(x, function(y) !any(is.na(y)))] }
incomplete <- function(x) {names( x[,sapply(x, function(y) any(is.na(y)))] )
}

trtr.na.var  <- incomplete(myTraining)
trts.na.var  <- incomplete(myTesting)
myTraining <- complete(myTraining)
myTesting  <- complete(myTesting)

#Lets remove the columns that are not predictors
myTraining_fset <- myTraining[,8:length(myTraining)]

#lets clear the variables with near zero variance
vNearZero <- nearZeroVar(myTraining_fset, saveMetrics = TRUE)
#Checking the result to see that there are zero values left
vNearZero
##
##          freqRatio percentUnique zeroVar  nzv
## roll_belt      1.065934      8.76358696  FALSE FALSE
## pitch_belt     1.017544     13.83322011  FALSE FALSE
## yaw_belt       1.052632     14.55502717  FALSE FALSE
## total_accel_belt 1.081769      0.23777174  FALSE FALSE
## gyros_belt_x    1.087173      1.09544837  FALSE FALSE
## gyros_belt_y    1.150382      0.55197011  FALSE FALSE
## gyros_belt_z    1.088571      1.33322011  FALSE FALSE
## accel_belt_x    1.065359      1.30774457  FALSE FALSE
## accel_belt_y    1.150108      1.12941576  FALSE FALSE
## accel_belt_z    1.099222      2.42866848  FALSE FALSE
## magnet_belt_x   1.022624      2.53906250  FALSE FALSE
## magnet_belt_y   1.009709      2.39470109  FALSE FALSE
## magnet_belt_z   1.010638      3.59205163  FALSE FALSE
## roll_arm       52.897436     19.53974185  FALSE FALSE
## pitch_arm      76.444444     22.20618207  FALSE FALSE
## yaw_arm       31.738462     21.53532609  FALSE FALSE
## total_accel_arm 1.000000      0.55197011  FALSE FALSE
## gyros_arm_x     1.087542      5.30740489  FALSE FALSE
## gyros_arm_y     1.513072      3.07404891  FALSE FALSE
## gyros_arm_z     1.056604      1.94463315  FALSE FALSE
## accel_arm_x     1.138614      6.44531250  FALSE FALSE
## accel_arm_y     1.141732      4.41576087  FALSE FALSE
## accel_arm_z     1.118421      6.39436141  FALSE FALSE
## magnet_arm_x    1.037037     11.09035326  FALSE FALSE
## magnet_arm_y    1.054545      7.25203804  FALSE FALSE
## magnet_arm_z    1.115942     10.53838315  FALSE FALSE
## roll_dumbbell   1.025000     87.66983696  FALSE FALSE
## pitch_dumbbell  2.390244     85.60631793  FALSE FALSE
## yaw_dumbbell    1.123288     87.25373641  FALSE FALSE
## total_accel_dumbbell 1.075518      0.34816576  FALSE FALSE
## gyros_dumbbell_x 1.053521      1.95312500  FALSE FALSE
## gyros_dumbbell_y 1.309735      2.22486413  FALSE FALSE
## gyros_dumbbell_z 1.058997      1.63043478  FALSE FALSE
## accel_dumbbell_x 1.045918      3.42221467  FALSE FALSE

```

```
## accel_dumbbell_y      1.100671      3.83831522      FALSE FALSE
## accel_dumbbell_z      1.097222      3.40523098      FALSE FALSE
## magnet_dumbbell_x      1.155340      8.91644022      FALSE FALSE
## magnet_dumbbell_y      1.219048      6.93783967      FALSE FALSE
## magnet_dumbbell_z      1.072727      5.59612772      FALSE FALSE
## roll_forearm          12.392473     14.93716033      FALSE FALSE
## pitch_forearm          60.631579     20.95788043      FALSE FALSE
## yaw_forearm            15.151316     14.30027174      FALSE FALSE
## total_accel_forearm    1.202817      0.56046196      FALSE FALSE
## gyros_forearm_x        1.044728      2.30129076      FALSE FALSE
## gyros_forearm_y        1.021645      6.04619565      FALSE FALSE
## gyros_forearm_z        1.107143      2.41168478      FALSE FALSE
## accel_forearm_x        1.056604      6.58967391      FALSE FALSE
## accel_forearm_y        1.000000      8.18614130      FALSE FALSE
## accel_forearm_z        1.076923      4.64504076      FALSE FALSE
## magnet_forearm_x       1.037037     11.96501359      FALSE FALSE
## magnet_forearm_y       1.347826     15.20889946      FALSE FALSE
## magnet_forearm_z       1.081081     13.38315217      FALSE FALSE
## classe                 1.469065      0.04245924      FALSE FALSE
```

Model Selection

The selected model for this task will be Random Forest because it generates an internal unbiased estimate of the generalization error as the forest building progresses, Random Forest works well with a mixture of numerical and categorical features

```
#The model chosen is the random forest method
keep <- names(myTraining_fset)
#fit model- RANDOM FOREST
set.seed(1235)

modFit <- randomForest(classe~., data = myTraining_fset)
print(modFit)
##
## Call:
## randomForest(formula = classe ~ ., data = myTraining_fset)
##           Type of random forest: classification
##           Number of trees: 500
## No. of variables tried at each split: 7
##
##           OOB estimate of  error rate: 0.73%
## Confusion matrix:
##      A      B      C      D      E  class.error
## A 3345      2      0      0      1 0.0008960573
## B   19 2252      8      0      0 0.0118473014
## C    0   12 2034      8      0 0.0097370983
## D    0    0  27 1902      1 0.0145077720
## E    0    0    2    6 2157 0.0036951501
qplot(roll_belt, magnet_dumbbell_y, colour=classe, data=myTraining_fset)
```



```

predict1 <- predict(modFit, myTesting, type = "class")
confusionMatrix(myTesting$classe, predict1)
## Confusion Matrix and Statistics
##
##              Reference
## Prediction    A      B      C      D      E
##      A 2229      3      0      0      0
##      B      8 1509      1      0      0
##      C      0      15 1350      3      0
##      D      0      0      11 1273      2
##      E      0      0      1      1 1440
##
## Overall Statistics
##
##              Accuracy : 0.9943
##              95% CI : (0.9923, 0.9958)
##      No Information Rate : 0.2851
##      P-Value [Acc > NIR] : < 2.2e-16
##
##              Kappa : 0.9927
##      McNemar's Test P-Value : NA
##
## Statistics by Class:
##
##              Class: A Class: B Class: C Class: D Class: E
## Sensitivity      0.9964      0.9882      0.9905      0.9969      0.9986
## Specificity      0.9995      0.9986      0.9972      0.9980      0.9997
## Pos Pred Value   0.9987      0.9941      0.9868      0.9899      0.9986
## Neg Pred Value   0.9986      0.9972      0.9980      0.9994      0.9997
## Prevalence       0.2851      0.1946      0.1737      0.1628      0.1838
## Detection Rate   0.2841      0.1923      0.1721      0.1622      0.1835
## Detection Prevalence 0.2845      0.1935      0.1744      0.1639      0.1838
## Balanced Accuracy 0.9979      0.9934      0.9938      0.9974      0.9992
predict_trainingset <- predict(modFit, myTraining, type = "class")
confusionMatrix(myTraining$classe, predict_trainingset)
## Confusion Matrix and Statistics
##
##              Reference
## Prediction    A      B      C      D      E
##      A 3348      0      0      0      0
##      B      0 2279      0      0      0
##      C      0      0 2054      0      0
##      D      0      0      0 1930      0
##      E      0      0      0      0 2165
##
## Overall Statistics
##
##              Accuracy : 1
##              95% CI : (0.9997, 1)
##      No Information Rate : 0.2843
##      P-Value [Acc > NIR] : < 2.2e-16
##
##              Kappa : 1
##      McNemar's Test P-Value : NA
##
## Statistics by Class:
##

```

	Class: A	Class: B	Class: C	Class: D	Class: E
## Sensitivity	1.0000	1.0000	1.0000	1.0000	1.0000
## Specificity	1.0000	1.0000	1.0000	1.0000	1.0000
## Pos Pred Value	1.0000	1.0000	1.0000	1.0000	1.0000
## Neg Pred Value	1.0000	1.0000	1.0000	1.0000	1.0000
## Prevalence	0.2843	0.1935	0.1744	0.1639	0.1838
## Detection Rate	0.2843	0.1935	0.1744	0.1639	0.1838
## Detection Prevalence	0.2843	0.1935	0.1744	0.1639	0.1838
## Balanced Accuracy	1.0000	1.0000	1.0000	1.0000	1.0000

Conclusion

As we can see from the results for the Training set, the random forest method is the best fit model and has been selected for the test data to submit the final results for this assignment as it shows an 100% accuracy.

```

predict_testset <- predict(modFit, testing_set, type = "class")
print(predict_testset)
##  1  2  3  4  5  6  7  8  9 10 11 12 13 14 15 16 17 18 19 20
##  B  A  B  A  A  E  D  B  A  A  B  C  B  A  E  E  A  B  B  B
## Levels: A B C D E

```